

Master in Artificial Intelligence

Advanced Human Language Technologies

Neural
Networks DDI

General
Structure

Detailed
Structure

Core task

Goals &
Deliverables



UNIVERSITAT POLITÈCNICA DE CATALUNYA
BARCELONATECH

Facultat d'Informàtica de Barcelona



Outline

1 Neural Networks DDI

2 General Structure

3 Detailed Structure

- Learner
- Classifier
- Auxiliary classes

4 Core task

5 Goals & Deliverables

Neural
Networks DDI

General
Structure

Detailed
Structure

Core task

Goals &
Deliverables

Session 6 - DDI using neural networks

Assignment

Write a python program that parses all XML files in the folder given as argument and recognizes and classifies sentences stating drug-drug interactions. The program must use a neural network approach.

```
$ python3 ./nn-DDI.py data/Devel/  
DDI-DrugBank.d398.s0|DDI-DrugBank.d398.s0.e0|DDI-DrugBank.d398.s0.e1|effect  
DDI-DrugBank.d398.s0|DDI-DrugBank.d398.s0.e0|DDI-DrugBank.d398.s0.e2|effect  
DDI-DrugBank.d211.s2|DDI-DrugBank.d211.s2.e0|DDI-DrugBank.d211.s2.e5|mechanism  
...
```

Neural
Networks DDI

General
Structure

Detailed
Structure

Core task

Goals &
Deliverables

Outline

1 Neural Networks DDI

2 General Structure

3 Detailed Structure

- Learner
- Classifier
- Auxiliary classes

4 Core task

5 Goals & Deliverables

Neural
Networks DDI

General
Structure

Detailed
Structure

Core task

Goals &
Deliverables

General Structure

The general structure is basically the same than for the traditional ML approach:

- Two programs: one learner and one classifier.
- The learner loads the training (Train) and validation (Devel) data, formats/encodes it appropriately, and feeds it to the model, together with the ground truth.
- The classifier loads the test data, formats/encodes it in the same way that was used in training, and feeds it to the model to get a prediction.

In the case of NN, we don't need to extract features (though we **do need** some encoding)

Input Encoding

- The input/output layers of a NN are vectors of neurons, each set to 0/1.
- Modern deep learning libraries handle this in the form of *indexes* (i.e. just provided the *position* of active neurons, omitting zeros).
- For instance, in a LSTM, each input word in the sequence may be encoded as the concatenation of different vectors each containing information about some aspect of the word (form, lemma, PoS, suffix...)
- Each vector will have only one active neuron, indicated by its *index*. This input is usually fed to an embedding layer.
- Our learned will need to create and store *index* dictionaries to be able to interpret the model later. See class *Codemaps* below.

Neural
Networks DDI

General
Structure

Detailed
Structure

Core task

Goals &
Deliverables

Outline

1 Neural Networks DDI

2 General Structure

3 Detailed Structure

- Learner
- Classifier
- Auxiliary classes

4 Core task

5 Goals & Deliverables

Neural
Networks DDI

General
Structure

Detailed
Structure

Core task

Goals &
Deliverables

Outline

1 Neural Networks DDI

2 General Structure

3 Detailed Structure

- Learner

- Classifier

- Auxiliary classes

4 Core task

5 Goals & Deliverables

Neural
Networks DDI

General
Structure

Detailed
Structure

Learner

Core task

Goals &
Deliverables

Learner - Main program

```
1 def learn(trainfile, validationfile, modelname) :
2     ## learns a NN model using trainfile as training data, and validationfile
3     ## as validation data. Saves learnt model in a file named modelname
4
5     # load train and validation data
6     traindata = Dataset(trainfile)
7     valdata = Dataset(validationfile)
8
9     # create indexes from training data
10    max_len = 150
11    codes = Codemaps(traindata, max_len, suf_len)
12
13    # build network
14    model = build_network(codes)
15    with redirect_stdout(sys.stderr) :
16        model.summary()
17
18    # encode datasets
19    Xt = codes.encode_words(traindata)
20    Yt = codes.encode_labels(traindata)
21    Xv = codes.encode_words(valdata)
22    Yv = codes.encode_labels(valdata)
23
24    # train model
25    with redirect_stdout(sys.stderr) :
26        model.fit(Xt, Yt, batch_size=32, epochs=10, validation_data=(Xv,Yv),
27                    verbose=1)
28
29    # save model and indexes
30    model.save(modelname)
31    codes.save(modelname)
```

Neural
Networks DDI

General
Structure

Detailed
Structure

Learner

Core task

Goals &
Deliverables

Outline

1 Neural Networks DDI

2 General Structure

3 Detailed Structure

- Learner
- Classifier
- Auxiliary classes

4 Core task

5 Goals & Deliverables

Neural
Networks DDI

General
Structure

Detailed
Structure

Classifier

Core task

Goals &
Deliverables

Classifier - Main program

```
1 def predict(modelname, datafile, outfile) :
2     '''
3     Loads a NN model from file 'modelname' and uses it to extract drugs
4     in datafile. Saves results to 'outfile' in the appropriate format.
5     '''
6
7     # load model and associated encoding data
8     model = load_model(fname)
9     codes = Codemaps(fname)
10
11    # load and encode data to annotate
12    testdata = Dataset(datafile)
13    X = codes.encode_words(testdata)
14
15    # tag sentences in dataset
16    Y = model.predict(X)
17    Y = [codes.idx2label(np.argmax(s)) for s in Y]
18
19    # extract relations
20    output_interactions(testdata, Y, outfile)
```

Note: Observe the output structure (one class per sentence+pair), different from the NER task (one class per token).

Neural
Networks DDI

General
Structure

Detailed
Structure

Classifier

Core task

Goals &
Deliverables

Outline

1 Neural Networks DDI

2 General Structure

3 Detailed Structure

- Learner

- Classifier

- **Auxiliary classes**

4 Core task

5 Goals & Deliverables

Neural
Networks DDI

General
Structure

Detailed
Structure

Auxiliary classes

Core task

Goals &
Deliverables

Auxiliary classes - parse_data

Processing the whole dataset with StanfordCore takes a long time, so it is convenient to run it once and for all:

```
1
2 # preprocess a dataset with StanfordCore, and store result in a
3 # pickle file for later use.
4 # usage: ./parse_data.py data-folder filename
5 #   e.g. ./parse_data.py ../../data/train train
6
7 datadir = sys.argv[1]
8 filename = sys.argv[2]
9
10 data = Dataset(datadir)
11 data.save(filename)
```

Neural
Networks DDI

General
Structure

Detailed
Structure

Auxiliary classes

Core task

Goals &
Deliverables

Auxiliary classes - Dataset

```
1 class Dataset:
2     ## constructor:
3     ## If 'filename' is a directory, parses all XML files in datadir,
4     ## tokenizes
5     ## each sentence, and stores a list of sentence/pairs, each
6     ## of them as a parsed tree with masked target entities
7     ## tokens (word, start, end, gold_label), plus associate ground truth.
8     ## If 'filename' is a '.pck' file, load data set pickle file
9     def __init__(self, filename)
10
11     ## saves dataset to a pickle file (to avoid repeating parsing)
12     def save(self, filename)
13
14     ## iterator to get sentences in the dataset
15     def sentence(self)
16     , , ,
```

Class Dataset will *mask* the target entities in the input sentence:

Original sentence: *Exposure to oral ketamine is unaffected by itraconazole compounds but greatly increased by ticlopidine.*

Pair	Masked sentence
e0-e1	Exposure to oral DRUG1 is unaffected by DRUG2 but greatly increased by DRUG_OTHER.
e0-e2	Exposure to oral DRUG1 is unaffected by DRUG_OTHER but greatly increased by DRUG2.
e1-e2	Exposure to oral DRUG_OTHER is unaffected by DRUG1 but greatly increased by DRUG2.

Neural
Networks DDI

General
Structure

Detailed
Structure

Auxiliary classes

Core task

Goals &
Deliverables

Auxiliary classes - Codemaps

```
1 class Codemaps :
2     # Constructor: create code mapper either from training data, or
3     #               loading codemaps from given file.
4     #               If 'data' is a Dataset, and lengths are not None,
5     #               create maps from given data.
6     #               If data is a string (file name), load maps from file.
7     def __init__(self, data, maxlen=None, suflen=None)
8     # Save created codemaps in file named 'name'
9     def save(self, name)
10    # Convert a Dataset into lists of word codes and suffix codes
11    # Adds padding and unknown word codes.
12    def encode_words(self, data)
13    # Convert the gold labels in given Dataset into a list of label codes.
14    # Adds padding
15    def encode_labels(self, data)
16    # get word index size
17    def get_n_words(self)
18    # get suf index size
19    def get_n_sufs(self)
20    # get label index size
21    def get_n_labels(self)
22    # get index for given word
23    def word2idx(self, w)
24    # get index for given suffix
25    def suff2idx(self, s)
26    # get index for given label
27    def label2idx(self, l)
28    # get label name for given index
29    def idx2label(self, i)
```

Neural
Networks DDI

General
Structure

Detailed
Structure

Auxiliary classes

Core task

Goals &
Deliverables

Required functions - build_network

```
1 def build_network(codes) :
2
3     # sizes
4     n_words = codes.get_n_words()
5     max_len = codes.maxlen
6     n_labels = codes.get_n_labels()
7
8     # word input layer & embeddings
9     inptW = Input(shape=(max_len,))
10    embW = Embedding(input_dim=n_words, output_dim=100,
11                    input_length=max_len, mask_zero=False)(inptW)
12
13    conv = Conv1D(filters=30, kernel_size=2, strides=1,
14                activation='relu', padding='same')(embW)
15    flat = Flatten()(conv)
16
17    out = Dense(n_labels, activation='softmax')(flat)
18
19    model = Model(inptW, out)
20    model.compile(loss='categorical_crossentropy', optimizer='adam',
21                metrics=['accuracy'])
22
23    return model
```

Neural
Networks DDI

General
Structure

Detailed
Structure

Auxiliary classes

Core task

Goals &
Deliverables

Required functions - build_network

- DDI is **not** a sequence tagging task (which assign one label per word), but a sentence classification, where a single label is assigned to the whole sentence (or sentence + entity pair in this case).
- Good results may be achieved using a CNN, as in the provided example.
- The problem also may be approached with an LSTM. Note that instead of getting the output at each word (`return_sequences=True`), only the output at the end of the sequence must be used (`return_sequences=False`).
- It is also possible to combine LSTM and CNN layers.
- You will need to add one Embedding layer after the input, that is where the created indexes will become handy.
- You may get inspiration for an architecture from these examples: [1], [2],[3],[4], some of the papers provided in labAHLT package in `papers/SharedTask/otherSystems`, or just googling for *semeval DDI neural networks*.

Neural
Networks DDI

General
Structure

Detailed
Structure

Auxiliary classes

Core task

Goals &
Deliverables

Outline

1 Neural Networks DDI

2 General Structure

3 Detailed Structure

- Learner
- Classifier
- Auxiliary classes

4 Core task

5 Goals & Deliverables

Neural
Networks DDI

General
Structure

Detailed
Structure

Core task

Goals &
Deliverables

Build a good NN-based DDI detector

Strategy: Experiment with different NN architectures and possibilities.

Some elements you can play with:

- Embedding dimensions, number and kind of layers, used optimizer...
- Using just CNN, just a LSTM, or a LSTM+CNN combination
- Using lowercased and/or non lowercased word embeddings
- Initializing embeddings with available pretrained model
- Using extra input (e.g. lemma embeddings, PoS embeddings, suffix/prefix embeddings, ...)
- Adding extra dense layers, with different activation functions
- Using pretrained transformers such as Bert as the first layers of your network.
- Adding attention layers
- ...etc.

Build a good NN-based DDI detector

Warnings:

- Neural Network training uses randomization, so different runs of the same program will produce different results. Run the same model several times.
- During training, Keras reports accuracy on training set and on validation set. Those values are usually over 85%. However, this is due to the fact that about 85% of the pairs have interaction “null” (no-interaction). Thus, 85% accuracies correspond very low F_1 values. To get a reasonable F_1 , val_accuracy must reach about ~89-90%.

To precisely evaluate how your model is doing, do not rely on reported accuracy: run the classifier on the Development set and use the evaluator.

Outline

1 Neural Networks DDI

2 General Structure

3 Detailed Structure

- Learner
- Classifier
- Auxiliary classes

4 Core task

5 Goals & Deliverables

Neural
Networks DDI

General
Structure

Detailed
Structure

Core task

Goals &
Deliverables

Exercise Goals

What you should do:

- Work on your architecture and input vectors. It is the component of the process where you have most control.
- Experiment with different architectures and hyperparameters.
- Experiment with different input information
- Keep track of tried variants and parameter combinations.

What you should **NOT** do:

- Alter the suggested code structure.
- Produce an overfitted model: If performance on the test dataset is much lower than on devel dataset, you probably are overfitting your model. Note that a very large/deep NN applied to a reduced dataset (as is this case) will just *remember* the training data and thus will not generalize properly.

Exercise Goals

Orientative results:

- Provided CNN architecture gets a macroaverage F1 over 50%. Input information includes only embeddings for word forms.
- The NN may be extended with extra input information, additional convolutional layers (either separate for each input or after concatenating them), changing their size, changing the size/stride of the convolutional kernel, adding LSTM layers (before the CNN, after it, or instead of it), maxpool layers (typically after the CNN or LSTM), etc.
- Goal: achieve ~65% macro average F1.

Deliverables

Write a report describing the work carried out in both NN exercises
The report must be a **single self-contained PDF document**, under ~10 pages, containing:

- *Introduction*: What is this report about. What is the goal of the presented work.
- *NN-based NERC*
 - *Architecture*: What architectures did you try, and which was finally selected.
 - *Input information*: What input data did you use, and how did you encode it to feed the NN.
 - *Code*: Include your `build_network` function (and **any other function** it may call), properly formatted and commented. **Do not include any other code..**
 - *Experiments and results*: Results obtained on the **devel** and **test** datasets, for different architecture/hyperparameter combinations you deem relevant.

Neural
Networks DDI

General
Structure

Detailed
Structure

Core task

Goals &
Deliverables

Deliverables (continued)

- *NN-based DDI*

- *Architecture*: What architectures did you try, and which was finally selected.
- *Input information*: What input data did you use, and how did you encode it to feed the NN.
- *Code*: Include your `build_network` function (and *any other function* it may call), properly formatted and commented. **Do not include any other code..**
- *Experiments and results*: Results obtained on the **devel** and **test** datasets, for different architecture/hyperparameter combinations you deem relevant.

- *Conclusions*: Final remarks and insights gained in this task.

Keep result tables in your report in the format produced by the evaluator module. Do not reorganize/summarize/reformat the tables or their content.

Neural
Networks DDI

General
Structure

Detailed
Structure

Core task

Goals &
Deliverables